
Spring 2022 ECE 6913 Section A, Quiz 2 Solutions

Problem 1. This problem explores energy efficiency and its relationship with performance. The parts of this problem assume the following energy consumption for activity in Instruction memory, Registers, and Data memory. You can assume that the other components of the datapath consume a negligible amount of energy. (“Register Read” and “Register Write” refer to the register file only.)

I-Mem	1 Register Read	Register Write	D-Mem Read	D-Mem Write
140pJ	70pJ	60pJ	140pJ	120pJ

Assume that components in the datapath have the following latencies. You can assume that the other components of the datapath have negligible latencies.

I-Mem	Control	Register Read or Write	ALU	D-Mem Read or Write
200 ps	150 ps	90 ps	90 ps	250 ps

1.1 How much energy is spent to execute an **addi** instruction in a single-cycle design and in the five-stage pipelined design

I-Mem is read, two registers are read, and a register is written

We have: $140\text{pJ} + 2 \times 70\text{pJ} + 60\text{pJ} = 340\text{pJ}$

1.2 How much energy is spent to execute a **lw** instruction in a single-cycle design

$140\text{pJ} + 2 \times 70\text{pJ} + 60\text{pJ} + 140\text{pJ} = 480\text{pJ}$

1.3 How much energy is spent to execute a **beq** instruction in a single-cycle design

$\text{I-Mem} + 2 \text{ registers} = 140\text{pJ} + 2 \times 70\text{pJ} = 280\text{pJ}$

Problem 2.

One difference between a write-through cache and a write-back cache can be in the time it takes to write. During the first cycle, we detect whether a hit will occur, and during the second (assuming a hit) we actually write the data.

Let's assume that 50% of the blocks are dirty for a write-back cache. For this question, assume that the write buffer for the write through will never stall the CPU (no penalty). Assume a cache read hit takes 1 clock cycle, the cache miss penalty is 50 clock cycles, and a block write from the cache to main memory takes 50 clock cycles. Finally, assume the instruction cache miss rate is 0.5% and the data cache miss rate is 1%. Assuming that on average 26% and 9% of instructions in the workload are loads and stores, respectively, **estimate the performance of a write-through cache with a two-cycle write versus a write-back cache with a two-cycle write.**

CPU performance equation: $CPUTime = IC * CPI * ClockTime$

$CPI = CPI_{execution} + StallCyclesPerInstruction$

We know:

Instruction miss penalty is 50 cycles

Data read hit takes 1 cycle

Data write hit takes 2 cycles

Data miss penalty is 50 cycles for write through cache

Data miss penalty is 50 cycles or 100 cycles for write back cache

Miss rate is 1% for data cache (MRD) and 0.5% for instruction cache (MRI)

50% of cache blocks are dirty in the write back cache

26% of all instructions are loads

9% of all instructions are stores

Then: $CPI_{execution} = 0.26 * 1 + 0.09 * 2 + 0.65 * 1 = 1.09$

Write through

$StallCyclesPerInstruction = MRI * 50 + MRD * (0.26 * 50 + 0.09 * 50) = 0.425$

so: $CPI = 1.09 + 0.425 = 1.515$ (1)

Write back

$StallCyclesPerInstruction = MRI * 50 + MRD * (0.26 * (0.5 * 50 + 0.5 * 100) + 0.09 * (0.5 * 50 + 0.5 * 100)) = 0.5125$

so: $CPI = 1.09 + 0.5125 = 1.6025$ (2)

Comparing 1 and 2 we notice that the system with the write back cache is 6% slower.

Problem 3.

Consider the following RISC V Instruction sequence executing in a 5-stage pipeline:

```
or    x13, x12, x11
ld    x10, 0(x13)
ld    x11, 8(x13)
add   x12, x10, x11
subi  x13, x12, 16
```

3.1 Identify all of the data hazards and their resolution with NOPs assuming no forwarding or hazard detection hardware is being used

Hazards identified:

```
or    x13, x12, x11
ld    x10, 0(x13)      EX to 1st RAW Hazard
ld    x11, 8(x13)      EX to 2nd RAW Hazard
add   x12, x10, x11    MEM to 1st RAW [load-use-data] & MEM to 2nd Hazards
subi  x13, x12, 16      Ex to 1st RAW Hazard
```

NOPS introduced to resolve Hazards:

```
or    x13, x12, x11
NOPS
NOPS
ld    x10, 0(x13)      EX to 1st RAW Hazard resolution with 2 NOPs
ld    x11, 8(x13)      EX to 2nd RAW Hazard resolved as well from above 2 NOPs
NOPS
NOPS
add   x12, x10, x11    MEM to 1st RAW [load-use-data] & MEM to 2nd Hazards
                        resolved with 2 NOPs
NOPS
NOPS
subi  x13, x12, 16      Ex to 1st only RAW Hazard resolved with 2 NOPs
```

3.2 If there is forwarding, for the first seven cycles during the execution of this code, *specify which signals are asserted in each cycle by hazard detection and forwarding units* in Figure below.

Mux control	Source	Explanation
ForwardA = 00	ID/EX	The first ALU operand comes from the register file.
ForwardA = 10	EX/MEM	The first ALU operand is forwarded from the prior ALU result.
ForwardA = 01	MEM/WB	The first ALU operand is forwarded from data memory or an earlier ALU result.
ForwardB = 00	ID/EX	The second ALU operand comes from the register file.
ForwardB = 10	EX/MEM	The second ALU operand is forwarded from the prior ALU result.
ForwardB = 01	MEM/WB	The second ALU operand is forwarded from data memory or an earlier ALU result.

	Clock Cycle	1	2	3	4	5	6	7	8	9	10
1	or	IF	ID	EX	MEM	WB					
2	ld		IF	ID	EX	MEM	WB				
3	ld			IF	ID	EX	MEM	WB			
4	NOP	mandatory NOP for which no forwarding solution possible: load-data-use									
5	add					IF	ID	EX	MEM	WB	
6	subi						IF	ID	EX	MEM	WB

- (1) A=x B=x (no instruction in EX stage yet)
- (2) A=x B=x (no instruction in EX stage yet)
- (3) A=0 B=0 (both operands of the or instruction: x11, x12 come from Reg File)
- (4) A=2 B=0 (base (RS1) in first ld (x13) taken from EX/MEM of previous instruction)
- (5) A=1 B=0 (base (RS1) in 2nd ld (x13) taken from MEM/WB of a previous instruction)
- (6) A=x B=x (no instruction in EX stage yet because NOP introduced to resolve MEM to 1st
- (7) A=0 B=1 (RS2 in the add instruction is x11 which is forwarded from MEM/WB of 2nd ld, the result of the 1st ld (x10) has already been written into Reg File in CC 6 - so, no forwarding necessary for first operand)
- (8) A=1 B=0 (RS1 of subi instruction forwarded from EX/MEM of add instruction)

Problem 4.

. Consider the following program and cache behaviors.

Data Reads per 1K instructions	Data Writes per 1K instructions	Instruction Cache Miss Rate	Data Cache Miss Rate	Block Size (Bytes)
300	150	0.5%	5%	128

Suppose a CPU with a write-through, write allocate cache achieves a CPI of 2.

4.1 What are the read and write bandwidths (measured by bytes per cycle) between RAM and the cache? (Assume each miss generates a request for one block.). *For a write-allocate policy, a write miss also makes a read request to RAM – please be sure to consider its impact on Read Bandwidth*

Instruction Bandwidth:

When the CPI is 2, there are, on average, 0.5 instruction accesses per cycle.

0.5 instructions read from Instruction memory per cycle

0.5% of these *instruction* accesses cause a cache **Read** miss (and subsequent memory request).

$$[0.5 \text{ instr/cycle}] \times [0.005 \text{ misses/instruction}] = \text{missed instructions/cycle}$$

Assuming each miss requests one block and each block is 128 bytes [16 words with 8 bytes (64 bits) per word] , instruction accesses generate an average of

$$[0.5 \text{ instr/cycle}] \times [0.005 \text{ misses/instruction}] \times [128 \text{ bytes/miss}] = \\ = 0.32 \text{ bytes/cycle of read traffic}$$

Read Data bandwidth:

30% of instructions generate a **read** request from data memory.

$$[0.5 \text{ instr/cycle}] \times [0.3 \text{ Read Data Accesses/instruction}] = [0.15 \text{ Read Data Accesses / cycle}]$$

5% of these generate a cache miss;

$$[0.15 \text{ Read Data Accesses / cycle}] \times [0.05 \text{ misses / Read Data Access}] = 0.0075 \text{ Read Misses/cycle}$$

Assuming each miss requests one block and each block is 128 bytes [16 words with 8 bytes (64 bits) per word] ,

$$[0.0075 \text{ Read Misses/cycle}] \times [128 \text{ Bytes/block}] \times [1 \text{ block/miss}] = 0.0075 \times 128 \text{ Bytes/cycle} \\ = 0.96 \text{ Bytes/cycle}$$

Write Data bandwidth:

15% of instructions generate a **write** request into data memory.

$$[0.5 \text{ instr/cycle}] \times [0.15 \text{ Write Data Accesses/instruction}] = [0.075 \text{ Write Data Accesses / cycle}]$$

All of the words written to the cache must be written into Memory:

$$[0.075 \text{ Write Data Accesses / cycle}] \times [8 \text{ bytes/word}] \times [1 \text{ word/write-through}] = 0.6 \text{ Bytes/cycle}$$

For a *Write-allocate policy*, a Write miss also makes a **read** request to RAM

$$[0.5 \text{ instr/cycle}] \times [0.15 \text{ Write Data Accesses/instruction}] \times [0.05 \text{ misses/Write Data Access}] \times [128 \text{ Bytes/miss}] \\ = 0.48 \text{ Bytes/cycle}$$

Assuming each miss requests one Word (8 bytes) since this is a write-through cache *with only 1 word written per miss into memory*,

$$[0.00375 \text{ Write Misses/cycle}] \times [8 \text{ Bytes/word}] \times [1 \text{ word/miss}] = 0.03 \text{ Bytes/cycle}$$

Total Read Bandwidth

$$0.32 \text{ (Instruction memory)} + 0.96 \text{ (data memory)} + 0.48 \text{ (Write-miss in Write-through cache with Write Allocate)} \text{ Bytes/cycle} = 1.76 \text{ Bytes/cycle}$$

Total Write Bandwidth:

$$0.6 \text{ Bytes/cycle} + 0.03 \text{ Bytes/cycle} = 0.63 \text{ Bytes/cycle}$$