

1. NVIDIA has a “half” format, which is similar to IEEE 754 except that it is only 16 bits wide. The leftmost bit is still the sign bit, the exponent is 5 bits wide and the Fraction field is 10 bits long. A hidden 1 is assumed.

REPRESENTATION RANGE OF THE NVIDIA ‘HALF FORMAT’



For each of the following, write the *binary value* and the **corresponding decimal value** of the 16-bit floating point number that is the closest available representation of the requested number. If rounding is necessary use round-to-nearest. Give the decimal values either as whole numbers or fractions.

(i) The number ‘0’, Charge of an electron (-1.6×10^{-19}), smallest positive *normalized* number, smallest positive *denormalized* number, largest positive *normalized* number, largest positive *denormalized* number, distance b/w Earth and Neptune in inches: 171,072,000,000,000

Number	Binary	Decimal
0	00000 0000000000	0
Charge of an electron: -1.6×10^{-19} (C)	1 00000 0000000000	0
Smallest positive normalized number	0 00001 0000000000	$6.103515625 \times 10^{-5}$
Smallest positive denormalized number > 0	0 00000 0000000001	$5.960464477 \times 10^{-8}$
Largest positive denormalized number > 0	0 00000 1111111111	$\sim 6.103515625 \times 10^{-5}$
Largest positive number < infinity	0 11110 1111111111	= 65536
Average distance b/w proton and neutron in Hydrogen atom = 0.8751×10^{-15} m	0 00000 0000000000	0
Distance between Earth and Neptune in inches = 171,072,000,000,000	0 11111 0000000000	overflow

Smallest normalized number > 0: 0 00001 0000000000; bias = 15; val of exponent = $1 - \text{bias} = -14$, Significand = $1.000...0_2 = 1$, Value in decimal $\sim 1.00 \times 2^{-14} = 6.103515625 \times 10^{-5}$

Smallest positive denormalized number > 0: 0 00000 0000000001; value of exponent = -14 ; significand = $0.0000000001 = 1 \times 2^{-10}$; so smallest denormalized number = $1 \times 2^{-10} \times 2^{-14} = 2^{-24} = 5.960464477 \times 10^{-8}$

Largest possible denormalized number > 0: 0 00000 1111111111; value of exponent = $1 - 15 = -14$; significand = $0.1111...1 \sim 1.0$; so largest denormalized number = $1 \times 2^{-14} \sim 6.103515625 \times 10^{-5}$

Largest positive number < inf: 0: 0 11110 1111111111; value of exponent = 30 - 15 = 15; Significand = 1.11...1₂ = ~ 2; value in decimal ~ 2 x 2¹⁵ = 2¹⁶ = 65536

Overflow: represented by +infinity: 0 11111 0000000000

2. Implement in RISC V these line of code in C:

(i) `f = g - A[B[C[27]]]`

(ii) `f = g - A[C[10] + B[11]]`

(iii) `A[i] = 4B[4i-44] + 3C[32i+32]`

3. When parallelizing an application, the ideal speedup is speeding up by the number of processors. This is limited by two things: percentage of the application that can be parallelized and the cost of communication. Amdahl's law takes into account the former but not the latter.

(i) What is the speedup with N processors if 80% of the application is parallelizable, ignoring the cost of communication?

$$\text{Speed-up} = 1 / (0.2 + 0.8/N)$$

What is the speedup with 8 processors if, *for every processor added, the communication overhead is 0.5% of the original execution time*?

$$\text{Speed-up} = 1 / (0.2 + 8 \times 0.005 + 0.8/8) = 2.94$$

What is the speedup with 8 processors if, for every time the *number of processors is doubled*, the communication overhead is *increased by 0.5% of the original execution time*?

$$\text{Speed-up} = 1 / (0.2 + 3 \times 0.005 + 0.8/8) = 3.17$$

(ii) Write the general equation that solves this question: What is the number of processors with the highest speedup in an application in which P% of the original execution time is parallelizable, and, for every time the number of processors is doubled, the communication is increased by 0.5% of the original execution time?

$$\frac{d}{dN} \left(\frac{1}{(1 - P) + \log N \times 0.005 + P/N} \right) = 0$$